

Choosing the Number of Clusters (K)

One of the most common questions in K-means is: "**How do I choose the value of K ?**"

Because clustering is unsupervised, there is often no single "right" answer. The data itself might be ambiguous—one person might see 2 clusters, while another sees 4.

However, there are two main methods to help you decide.

1. The Elbow Method (Academic Approach)

This is a technique to automatically visualize the "best" K , though it is not always reliable.

- **How it works:**

1. Run K-means with a variety of K values (e.g., $K = 1, 2, 3, \dots, 10$).
2. For each K , calculate the final **Cost Function** (Distortion) J .
3. Plot the cost J (y-axis) against the number of clusters K (x-axis).

- **What to look for:**

- As K increases, the cost J will always go down.
- You are looking for an "**Elbow**" in the curve—a point where the cost decreases rapidly and then suddenly flattens out. That point is your optimal K .

- **The Limitation:**

- In real-world applications, the curve is often very smooth with no distinct "elbow." This makes the method ambiguous and difficult to use reliably.
- **Important:** You should *never* choose K simply to minimize the cost function J . If you did, you would always choose the largest possible K (where every point is its own cluster), which is useless.

2. The "Downstream Purpose" Method (Recommended)

Since you are usually clustering data to solve a specific real-world problem, the best way to choose K is to ask: "**Which K helps me solve my problem best?**"

- **Example: T-Shirt Sizing**

- Imagine you have height and weight data for customers.
- **Option A ($K = 3$):** You create Small, Medium, and Large sizes. This is cheaper to manufacture and market.
- **Option B ($K = 5$):** You create XS, S, M, L, XL. This fits customers better but is more expensive to manufacture.
- **Decision:** You run K-means for both $K = 3$ and $K = 5$, then compare the business

trade-offs (Customer Fit vs. Manufacturing Cost) to decide.

- **Example: Image Compression**

- You can use K-means to compress an image by reducing the number of colors.
 - **Trade-off:** You balance **Image Quality** (higher K) vs. **File Size** (lower K).
-

Summary of K-Means

Congratulations! You have finished the section on K-means clustering. You now know:

- How the algorithm works (Assign points → Move centroids).
- The optimization objective (Cost function J).
- How to initialize it (Random initialization).
- How to choose K (Downstream purpose).